



0x Permit2Payment

Security Assessment (Summary Report)

June 18, 2024

Prepared for:

Eric Wong and Duncan Townsend

ZeroEx

Prepared by: **Elvis Skoždopolj and Kurt Willis**

About Trail of Bits

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at info@trailofbits.com.

Trail of Bits, Inc.

497 Carroll St., Space 71, Seventh Floor
Brooklyn, NY 11215

<https://www.trailofbits.com>

info@trailofbits.com

Notices and Remarks

Copyright and Distribution

© 2024 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be public information; it is licensed to ZeroEx under the terms of the project statement of work and has been made public at ZeroEx's request. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

Table of Contents

About Trail of Bits	1
Notices and Remarks	2
Table of Contents	3
Project Summary	4
Project Targets	5
Executive Summary	6
Codebase Maturity Evaluation	8
Summary of Findings	10
1. Internal function pointer masking could become insufficient in the future	11
A. Vulnerability Categories	12
B. Code Maturity Categories	14
C. Code Quality Issues	16
D. Mutation Testing	18
Unit Tests	20
Integration Tests	22
Recommendations	23

Project Summary

Contact Information

The following project manager was associated with this project:

Jeff Braswell, Project Manager
jeff.braswell@trailofbits.com

The following engineering director was associated with this project:

Josselin Feist, Engineering Director, Blockchain
josselin.feist@trailofbits.com

The following consultants were associated with this project:

Elvis Skoždopolj, Consultant
elvis.skozdopolj@trailofbits.com

Kurt Willis, Consultant
kurt.willis@trailofbits.com

Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
May 29, 2024	Pre-project kickoff call
June 10, 2024	Delivery of report draft
June 10, 2024	Report readout meeting
June 18, 2024	Delivery of summary report

Project Targets

The engagement involved a review and testing of the following target.

Permit2Payment

Repository	https://github.com/0xProject/0x-settler
Version	f6e15f6
Type	Solidity
Platform	Ethereum

Executive Summary

Engagement Overview

ZeroEx engaged Trail of Bits to review the security of the **Permit2Payment** smart contracts. These contracts are meant to be inherited by the main contracts, **Settler** and **SettlerMetaTx**, which enable execution of optimized and configurable batched actions. The **Permit2Payment** contracts implement internal functions and libraries that manage how transient storage values are set, checked, and cleared, how calls to the **Permit2** and **AllowanceHolder** contracts are made, and how callbacks are validated and executed. The main goal of the contracts is to securely handle token allowances and permits via calls and callbacks performed through the 0x Settler protocol.

Two consultants conducted the review from May 30 to June 7, 2024, for a total of two engineer-weeks of effort. With full access to source code and documentation, we performed static and dynamic testing of the target, using automated and manual processes.

Observations and Impact

The review was focused on checking whether the **Permit2Payment** contracts can correctly and securely handle calls related to token transfers and callbacks performed through the 0x Settler protocol or any valid target of the protocol, with a focus on the correct handling of **Permit2** calls. We reviewed how transient storage slots are set, cleared, and validated in order to ensure that “old” values cannot be misused, that the slots cannot be arbitrarily modified, and that the risk of reentrancy is properly mitigated. Additionally, we checked that the assembly is correct, that masking cannot truncate important values such as the internal function pointers, that values are correctly packed inside the operator slot, and that arbitrary callbacks cannot be performed.

Interactions with other contracts in the codebase were reviewed to gain context on the use of **Permit2Payment**; however, they were not the main focus of the review. Findings listed in the previous 0x Settler security review report were omitted.

During the review, we did not identify any findings that currently have a direct impact on the security of the codebase; however, we did identify one informational-severity issue (**TOB-0XP2-1**) that should be addressed prior to deploying the contracts on new chains. Additionally, the codebase uses advanced low-level optimization techniques and new Solidity opcode features, minimizing unnecessary checks and optimizing gas costs. While these techniques and features are currently implemented securely, they do make the system more difficult to review and maintain. They also pose a risk that minor and difficult-to-find mistakes and as-yet undiscovered compiler bugs could impact the security of the protocol.

Recommendations

We recommend that the ZeroEx team address the finding disclosed in this report prior to deploying the contracts on new chains. Also, due to the heavy use of assembly and new Solidity opcodes, we recommend that the ZeroEx team simplify the assembly-heavy code, document all unwritten assumptions (e.g., the internal function pointer size assumption), and regularly check for compiler bugs that could impact the features being used in the codebase.

Codebase Maturity Evaluation

Trail of Bits uses a traffic-light protocol to provide each client with a clear understanding of the areas in which its codebase is mature, immature, or underdeveloped. Deficiencies identified here often stem from root causes within the software development life cycle that should be addressed through standardization measures (e.g., the use of common libraries, functions, or frameworks) or training and awareness programs.

Category	Summary	Result
Arithmetic	The target contracts do not use any arithmetic operations.	Not Applicable
Auditing	The target contracts do not emit any events; events are emitted by the inheriting contracts.	Not Applicable
Authentication / Access Controls	The target contracts implement access controls and reentrancy protection based on the values of transient storage slots. The <code>Permit2</code> interactions on behalf of users other than the caller are performed via the <code>SettlerMetaTx</code> contract and are protected by the witness parameter calculated in <code>SettlerMetaTx</code> . Arbitrary callbacks are prevented by the setting of the function selector, internal function pointer, and expected caller in a transient storage slot, and by the tight control over the callback data.	Satisfactory
Complexity Management	<p>The callback mechanism uses transient storage slots, packed values, and internal function pointers to determine which callback function will be triggered and provide access controls. While this provides additional flexibility, it also creates overhead during code reviews and could be error-prone.</p> <p>While the project extensively uses inline assembly and heavy optimizations, the <code>Permit2Payment</code> contracts have a relatively simple design compared to the rest of the system.</p>	Moderate
Decentralization	The target contracts and the protocol are permissionless, and no privileged roles or actions exist. The contracts are not upgradeable and have no configuration parameters.	Strong

Documentation	The project documentation includes diagrams that detail various user flows and expected contract interactions. The inline documentation is comprehensive; however, the documentation of some assembly blocks and inherent system assumptions could be improved.	Satisfactory
Low-Level Manipulation	The Permit2Payment contracts use low-level manipulation for setting, fetching, and clearing transient storage slots, as well as for packing different values into a single slot. The system is designed with the goal of minimizing gas consumption; some checks are skipped if they appear elsewhere in the code or are considered unnecessary. However, the use of new Solidity features and low-level manipulation increases the chance that a small mistake made during future development could have significant security implications.	Moderate
Testing and Verification	The testing suite heavily uses mock calls, which lowers the effectiveness of the tests and could prevent issues from being discovered. Mutation testing indicates that gaps exist in the coverage of the testing suite (for more information, see appendix D).	Moderate
Transaction Ordering	We did not uncover any front-running, back-running, or sandwich-related issues.	Satisfactory

Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	Internal function pointer masking could become insufficient in the future	Data Validation	Informational

1. Internal function pointer masking could become insufficient in the future

Severity: Informational

Difficulty: High

Type: Data Validation

Finding ID: TOB-0XP2-1

Target: `src/core/Permit2Payment.sol`

Description

When the internal function pointer value is packed into the operator transient storage slot, the callback variable's higher bits are cleaned by using a bitwise AND operation, preserving the least significant 2 bytes of data.

```
function setOperatorAndCallback(
    address operator,
    uint32 selector,
    function (bytes calldata) internal returns (bytes memory) callback
) internal {
    // ...
    assembly ("memory-safe") {
        tstore(
            _OPERATOR_SLOT,
            or(
                shl(0xe0, selector),
                or(shl(0xa0, and(0xffff, callback)),
            and(0xffffffffffffffffffffffffffffffff, operator))
            )
        )
    }
}
```

Figure 1.1: Packing of the callback pointer into the operator transient storage slot
(`0x-settler/src/core/Permit2Payment.sol#57-64`)

Internal function pointers are encoded as jump locations in the contract code. **EIP-170 (Spurious Dragon)** introduces a maximum contract code size limit of 24 KB (0x6000 in hexadecimal); this means that, for now, 2 bytes are enough to store any internal function pointer on Ethereum. However, this limit could change in the future. Also, other networks might opt for a higher contract size limit. If the code is deployed to a network with a higher contract size limit, this code might not perform as expected.

Recommendations

Short term, consider using 4 bytes of data for the internal function pointer.

Long term, document the assumptions around the byte masking for internal function pointers and ensure they are checked for each deployment to a new chain.

A. Vulnerability Categories

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

B. Code Maturity Categories

The following tables describe the code maturity categories and rating criteria used in this document.

Code Maturity Categories	
Category	Description
Arithmetic	The proper use of mathematical operations and semantics
Auditing	The use of event auditing and logging to support monitoring
Authentication / Access Controls	The use of robust access controls to handle identification and authorization and to ensure safe interactions with the system
Complexity Management	The presence of clear structures designed to manage system complexity, including the separation of system logic into clearly defined functions
Cryptography and Key Management	The safe use of cryptographic primitives and functions, along with the presence of robust mechanisms for key generation and distribution
Decentralization	The presence of a decentralized governance structure for mitigating insider threats and managing risks posed by contract upgrades
Documentation	The presence of comprehensive and readable codebase documentation
Low-Level Manipulation	The justified use of inline assembly and low-level calls
Testing and Verification	The presence of robust testing procedures (e.g., unit tests, integration tests, and verification methods) and sufficient test coverage
Transaction Ordering	The system's resistance to transaction-ordering attacks

Rating Criteria	
Rating	Description
Strong	No issues were found, and the system exceeds industry standards.
Satisfactory	Minor issues were found, but the system is compliant with best practices.
Moderate	Some issues that may affect system safety were found.

Weak	Many issues that affect system safety were found.
Missing	A required component is missing, significantly affecting system safety.
Not Applicable	The category is not applicable to this review.
Not Considered	The category was not considered in this review.
Further Investigation Required	Further investigation is required to reach a meaningful conclusion.

C. Code Quality Issues

The following list highlights areas where the repository's code quality could be improved.

- The `setOperatorAndCallback` and `setWitness` function of the `TransientStorage` library do not check that the inputs are not zero. While these parameters in their current use cannot be zero due to their derivation, the assumptions should be clearly documented.

```
function setOperatorAndCallback(
    address operator,
    uint32 selector,
    function (bytes calldata) internal returns (bytes memory) callback
) internal {
    // ...

    assembly ("memory-safe") {
        tstore(
            _OPERATOR_SLOT,
            or(
                shl(0xe0, selector),
                or(shl(0xa0, and(0xffff, callback)),
and(0xffffffffffffffffffffffffffffffff, operator))
            )
        )
    }
}
```

Figure C.1: The `setOperatorAndCallback` function does not perform a zero-value check. ([0x-settler/src/core/Permit2Payment.sol#35-66](#))

- The constants `GENERIC`, `ASSERT_FAIL`, `CORRUPT_STORAGE_ARRAY`, `POP_EMPTY_ARRAY`, and `ZERO_FUNCTION_POINTER` from the `Panic` library are never used.

```
library Panic {
    // ...

    uint8 internal constant GENERIC = 0x00;
    uint8 internal constant ASSERT_FAIL = 0x01;
    uint8 internal constant ARITHMETIC_OVERFLOW = 0x11;
    uint8 internal constant DIVISION_BY_ZERO = 0x12;
    uint8 internal constant ENUM_CAST = 0x21;
    uint8 internal constant CORRUPT_STORAGE_ARRAY = 0x22;
    uint8 internal constant POP_EMPTY_ARRAY = 0x31;
    uint8 internal constant ARRAY_OUT_OF_BOUNDS = 0x32;
    uint8 internal constant OUT_OF_MEMORY = 0x41;
    uint8 internal constant ZERO_FUNCTION_POINTER = 0x51;
```

```
}
```

*Figure C.2: The Panic library defines unused error constants.
(0x-settler/src/utils/Panic.sol#4-24)*

- The `getAndClearWitness` function in the `TransientStorage` library uses decimal notation for zero, while the `clearPayer` and `getAndClearOperatorAndCallback` functions use hexadecimal notation.

```
function getAndClearWitness() internal returns (bytes32 witness) {  
    assembly ("memory-safe") {  
        witness := tload(_WITNESS_SLOT)  
        tstore(_WITNESS_SLOT, 0)  
    }  
}
```

*Figure C.3: The `getAndClearWitness` function
(0x-settler/src/core/Permit2Payment.sol#115-120)*

```
function clearPayer(address expectedOldPayer) internal {  
    address oldPayer;  
    assembly ("memory-safe") {  
        oldPayer := tload(_PAYER_SLOT)  
    }  
    if (oldPayer != expectedOldPayer) {  
        revert PayerSpent();  
    }  
    assembly ("memory-safe") {  
        tstore(_PAYER_SLOT, 0x00)  
    }  
}
```

*Figure C.4: The `clearPayer` function
(0x-settler/src/core/Permit2Payment.sol#144-155)*

D. Mutation Testing

This appendix outlines how we conducted mutation testing and highlights some of the most actionable results.

At a high level, mutation tests make several changes to each line of a target file and rerun the test suite for each change. Changes that result in test failures indicate adequate test coverage, while changes that do not result in test failures indicate gaps in the test coverage. Although mutation testing is a slow process, it allows auditors to focus their review on areas of the codebase that are most likely to contain latent bugs, and it allows developers to identify and add missing tests.

We used an experimental new mutation tool, `slither-mutate`, to conduct our mutation testing campaign. This tool is custom-made for Solidity and features higher performance and fewer false positives than existing tools such as `universalmutator`.

```
python3 -m pip install slither-analyzer
```

Figure D.1: The command that installs slither using pip

The mutation campaign was run against the smart contracts using the following commands:

```
slither-mutate . --test-cmd='forge test' --test-dir='test'
--ignore-dirs='script,lib,test,node_modules,deployer,multicall,vendor' --timeout 120
```

Figure D.2: A Bash script that runs a mutation testing campaign against each Solidity file in the src directory (unit tests)

```
slither-mutate . --test-cmd='FOUNDRY_PROFILE=integration forge test'
--test-dir='test'
--ignore-dirs='script,lib,test,node_modules,deployer,multicall,vendor' --timeout 130
```

Figure D.3: A Bash script that runs a mutation testing campaign against each Solidity file in the src directory (integration tests)

Consider the following notes about the above commands:

- On a consumer-grade laptop, the overall runtime of the mutation testing campaign is approximately 11 hours.
- The `--test-cmd` flags specify the command to run to assess mutant validity. A `--fail-fast` or `--bail` flag will automatically be added to test commands to improve runtime.

An abbreviated, illustrative example of a mutation test output file is shown in figure D.4.

```

INFO:Slither-Mutate:Mutating contract UniswapV3Fork
INFO:Slither-Mutate:[RR] Line 121: 'Panic.panic(Panic.ARITHMETIC_OVERFLOW)' ==> 'revert()' -->
UNCAUGHT
INFO:Slither-Mutate:[RR] Line 134: '(token0, token1) = (token1, token0)' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 144: '_updateSwapCallbackData(swapCallbackData, token1, payer)' ==>
'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 158: 'tempPrice = MIN_PRICE_SQRT_RATIO + 1' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 160: 'tempPrice = MAX_PRICE_SQRT_RATIO - 1' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 162: '(amount0, amount1) = abi.decode(
    _setOperatorAndCall(
        address(pool),
        abi.encodeCall(
            pool.swap,
            (
                // Intermediate tokens go to this contract.
                address(this),
                zeroForOne,
                int256(sellAmount),
                tempPrice,
                swapCallbackData
            )
        ),
        uint32(callbackSelector),
        _uniV3ForkCallback
    ),
    (int256, int256)
)' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 189: 'tempPrice = MAX_PRICE_SQRT_RATIO - 1' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 216: '_buyAmount = -(amount0)' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 219: 'Panic.panic(Panic.ARITHMETIC_OVERFLOW)' ==> 'revert()' -->
UNCAUGHT
INFO:Slither-Mutate:[RR] Line 228: 'payer = address(this)' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 229: 'sellAmount = buyAmount' ==> 'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[RR] Line 231: 'encodedPath = _shiftHopFromPathInPlace(encodedPath)' ==>
'revert()' --> UNCAUGHT
INFO:Slither-Mutate:[CR] Line 154: 'freeMemPtr := mload(0x40)' ==> '//freeMemPtr := mload(0x40)' -->
UNCAUGHT
INFO:Slither-Mutate:[CR] Line 182: 'mstore(0x40, freeMemPtr)' ==> '//mstore(0x40, freeMemPtr)' -->
UNCAUGHT
INFO:Slither-Mutate:[CR] Line 233: 'mstore(swapCallbackData, SWAP_CALLBACK_PREFIX_DATA_SIZE)' ==>
'//mstore(swapCallbackData, SWAP_CALLBACK_PREFIX_DATA_SIZE)' --> UNCAUGHT
INFO:Slither-Mutate:[FHR] Line 112: 'function _uniV3ForkSwap(
    address recipient,
    bytes memory encodedPath,
    uint256 sellAmount,
    uint256 minBuyAmount,
    address payer,
    bytes memory swapCallbackData
) internal returns (uint256 buyAmount) ' ==> 'function _uniV3ForkSwap(
    address recipient,
    bytes memory encodedPath,
    uint256 sellAmount,
    uint256 minBuyAmount,
    address payer,
    bytes memory swapCallbackData
) private returns (uint256 buyAmount) ' --> UNCAUGHT
INFO:Slither-Mutate:[FHR] Line 242: 'function _isPathMultiHop(bytes memory encodedPath) private pure
returns (bool) ' ==> 'function _isPathMultiHop(bytes memory encodedPath) private view returns (bool)
' --> UNCAUGHT
...

```

Figure D.4: Abbreviated output from the mutation testing campaign assessing test coverage of the in-scope contracts

Unit Tests

The following table displays the portion of each type of mutant for which all unit tests passed. The presence of valid mutants indicates that there are gaps in test coverage because the test suite did not catch the introduced change.

Contracts supporting tests, contracts that produced zero analyzed mutants (e.g., interfaces), and contracts for which mutation testing failed were omitted from the mutation testing analysis.

We recommend running the commands in figures D.2 and D.3 to get a full list of uncaught mutants.

Target (Unit Tests)	Uncaught Reverts	Uncaught Comments	Uncaught Tweaks
SettlerMetaTx	100% (13/13)	100% (13/13)	None analyzed
Settler	0% (0/13)	0% (0/16)	81.6% (31/38)
CurveTricrypto	100% (6/6)	100% (26/26)	None analyzed
MakerPSM	0% (0/5)	20% (1/5)	14.3% (3/21)
UniswapV3Fork	50% (12/24)	23% (3/13)	46% (37/80)
Permit2Payment	14.3% (1/7)	0% (0/7)	33.33% (1/3)
RfqOrderSettlement	10% (1/10)	0% (0/16)	None analyzed

The following is a summary of the mutations that were not caught by the unit testing suite:

- SettlerMetaTx
 - Replacing line 29, 37, 48, 112, 121, 130, 132, 134, 143, 155, 162, or 163 with `revert()`

- Commenting out line 33, 64, 65, 69, 70, 72, 88, 89, 90, 91, 151, or 158
- **Settler**
 - Replacing line 90 with `false` or negating the condition (e.g., `!condition`)
 - Replacing line 144 with `true` or `false`
 - Replacing the operator `==` with `<`, `>`, `<=`, `>=`, or `!=` at line 69, 90, 112, or 114
- **CurveTricrypto**
 - Replacing line 88, 109, 121, 122, 128, or 166 with `revert()`
 - Commenting out line 71, 72, 73, 79, 80, 84, 85, 114, 116, 117, 118, 119, 134, 135, 136, 137, 138, 139, 140, 146, 147, 148, 152, 153, 155, or 156
- **MakerPSM**
 - Commenting out line 41
 - Updating the variable on line 36 to `immutable` or to the `int128` type
 - Updating the variable on line 38 to `immutable`
- **UniswapV3Fork**
 - Replacing line 121, 134, 144, 158, 160, 162, 189, 216, 219, 228, 229, or 231 with `revert()`
 - Commenting out line 154, 182, or 233
 - Replacing line 120, 141, 151, 157, 186, 213, 218, or 223 with `true`, `false`, or a negated condition
 - Updating the `>` operator on line 120 to `>=` or `==`
 - Updating the operator `<` on line 133 to `<=` or `!=`
 - Updating the operator `<` on line 218 to `<=` or `==`
 - Updating the operator `<` on line 236 to `!=`
 - Updating `uint256` to `uint128` on line 41, 44, 46, 48, 49, 50, 58, or 152
 - Updating line 120, `sellAmount > uint256(type(int256).max)`, to `sellAmount > uint128(type(int128).max)`
- **Permit2Payment**

- Replacing line 255 with `revert()`
 - Replacing line 242 with `false`
- `RfqOrderSettlement`
 - Replacing line 145 with `revert()`

Integration Tests

The following table displays the portion of each type of mutant for which all integration tests passed.

Target (Integration Tests)	Uncaught Reverts	Uncaught Comments	Uncaught Tweaks
SettlerMetaTx	15.4% (2/13)	39.4% (7/23)	None analyzed
SettlerAbstract	0% (0/1)	100% (0/1)	0% (0/2)
SettlerBase	16.6% (1/6)	80% (8/10)	0% (0/2)
Context	33.3% (1/3)	50% (1/2)	33.3% (1/3)
Settler	0% (0/13)	56.2% (9/16)	18.4% (7/38)
MakerPSM	80% (4/5)	100% (1/1)	90.4% (19/21)
CurveTricrypto	0% (0/6)	41.9% (13/31)	None analyzed

The following is a summary of the mutations that were not caught by the integration testing suite:

- `SettlerMetaTx`
 - Replacing line 37 or 132 with `revert()`

- Commenting out line 33, 143, 151, 158, 162, or 163
- **SettlerAbstract**
 - Commenting out line 14
- **SettlerBase**
 - Commenting out line 58, 59, 80, or 86
- **Context**
 - Replacing line 18 with `revert()`
 - Commenting out line 22
- **Settler**
 - Commenting out line 18, 31, 33, 65, 89, 99, 115, 120, or 162
 - Updating the operator `==` on line 24 to `<=` or `>=`
 - Updating the operator `==` on line 90 to `<=`
 - Updating the operator `==` on line 114 to `<=` or `>=`
- **MakerPSM**
 - Replacing line 47, 48, 59, or 60 with `revert()`
 - Replacing the operator `*` with `+`, `/`, `-`, or `%` on line 46 or 53
 - Replacing the operator `+` with `/`, `-`, `*`, or `%` on line 55
 - Removing the variable assignment on line 46, 53, 55, or 57
- **CurveTricrypto**
 - Commenting out line 71, 109, 114, 117, 119, 122, 128, 134, 135, 137, 139, 140, or 156

Recommendations

We recommend that the ZeroEx team review the existing tests and add additional verification to catch the aforementioned types of mutations. Then, use a script similar to the ones provided in figure D.2 and figure D.3 to rerun a mutation testing campaign to ensure that the added tests provide adequate coverage.